

HEMANTH KAVULA

Pitman, NJ | Email: Kavula.hemanth@gmail.com | Phone: 856-524-6050 |

LinkedIn: <https://www.linkedin.com/in/hemanthkavula/>

PROFESSIONAL SUMMARY

Senior Data Engineer with 5+ years of experience building scalable data pipelines and modern lakehouse data solutions across multi-cloud environments. Strong expertise in Python, SQL, Spark, and distributed data processing with hands-on experience in ETL development, data integration, data governance, analytics engineering, and real-time streaming pipelines. Known for delivering reliable, high-performance data workflows that support enterprise analytics, reporting, and large-scale data operations.

TECHNICAL SKILLS

Programming & Querying: Python, SQL, Spark SQL

Data Engineering: ETL Development, Data Pipelines, Data Integration, Data Modeling, Data Governance, Data Warehousing, Lakehouse Architecture

Big Data & Processing: Apache Spark, PySpark, Distributed Data Processing, Batch & Streaming Pipelines

Cloud Platforms: Microsoft Azure, Amazon Web Services (AWS), Google Cloud Platform (GCP)

Data Warehousing & Modeling: Snowflake, Azure Synapse Analytics, Dimensional Modeling, Star Schema

Streaming & Messaging: Kafka, Event Hubs, Real-Time Data Processing

ETL & Orchestration: Azure Data Factory, Apache Airflow, DBT, AWS Glue

Databases: SQL Server, PostgreSQL, Oracle, MySQL

DevOps & Monitoring: CI/CD, Terraform, Docker, Monitoring & Logging

Data Formats: Parquet, JSON, CSV, Delta

PROFESSIONAL EXPERIENCE

Senior Data Engineer — Fidelity Investments

Jersey City, NJ | Jan 2025 – Present

- Build scalable lakehouse data pipelines using Python, Spark, Azure data services, and AWS storage platforms to process large enterprise datasets.
- Develop automated ETL and data integration workflows supporting ingestion of structured, semi-structured, and streaming data sources.

- Create dimensional data models and curated data layers supporting analytics and reporting use cases.
- Implement real-time streaming pipelines enabling near real-time analytics and operational reporting.
- Optimize SQL queries and Spark transformations improving performance and processing efficiency.
- Maintain data governance and data quality practices including schema validation, monitoring, and secure access controls.
- Collaborate with analytics teams and engineering stakeholders to deliver reliable analytics-ready datasets.
- Support CI/CD processes and automate deployment of data workflows across Azure and AWS environments.
- Participate in Agile development practices including sprint planning and technical discussions.
- Troubleshoot production data pipelines and improve workflow stability through monitoring and optimization.

Data Engineer — Fiserv

Bangalore, India | Jan 2020 – Dec 2023

- Develop AWS-based ETL pipelines using Python, SQL, and Spark to process high-volume enterprise datasets.
- Build batch and streaming ingestion workflows leveraging AWS cloud storage and event-driven processing.
- Create dimensional data models improving reporting performance and analytics usability.
- Integrate APIs, relational databases, flat files, and cloud storage platforms enabling seamless data integration across systems.
- Optimize database queries and transformation logic improving pipeline efficiency and performance.
- Implement CI/CD deployment workflows ensuring consistent delivery of data solutions across environments.
- Migrate legacy ETL processes into orchestrated pipelines improving scalability and maintainability.
- Collaborate with analytics teams to deliver production-ready datasets aligned with business reporting needs.
- Maintain monitoring and validation processes ensuring data governance and operational reliability.
- Support performance tuning and debugging of large-scale data workflows.

EDUCATION

Master of Science in Computer Science — Rowan University, Glassboro, NJ | Jan 2024 – Dec 2025